

Adaboost

Aaditya Ramdas

Dept. of Statistics and Data Science
Machine Learning Dept.
Carnegie Mellon University

Outline

1. Weak learning implies strong learning ($1/2$ class)

2. *Adaboost* ($1/2$ class)

Recap

A “mixed” strategy is a distribution over actions.

Expected payoff is $\mathbb{E}_{r \sim p, c \sim q}[M(r, c)] = p^T M q = \sum_{r \in [R]} \sum_{c \in C} M(r, c) p_r q_c$

Theorem: $\min_{q \in \Delta_C} \max_{p \in \Delta_R} p^T M q = \max_{p \in \Delta_R} \min_{q \in \Delta_C} p^T M q = v^*$

Value of the game

e_j is the canonical basis vector $[0, \dots, 1, \dots, 0]$

Implications:

$$\begin{aligned} \exists p \in \Delta_R \forall q \in \Delta_C \quad p^T M q &\geq v^* \\ \forall q \in \Delta_C \exists i \in [R] \quad e_i^T M q &\geq v^* \\ \exists q \in \Delta_C \forall p \in \Delta_R \quad p^T M q &\leq v^* \\ \forall p \in \Delta_R \exists j \in [C] \quad p^T M e_j &\leq v^* \end{aligned}$$

For “mixed strategies”, order does not matter!

Zero-sum games [\[edit \]](#)

The minimax theorem was first proven and published in 1928 by [John von Neumann](#),^[3] who is quoted as saying "*As far as I can see, there could be no theory of games ... without that theorem ... I thought there was nothing worth publishing until the Minimax Theorem was proved*".^[4]

Formally, von Neumann's minimax theorem states:

Let $X \subset \mathbb{R}^n$ and $Y \subset \mathbb{R}^m$ be [compact convex](#) sets. If $f : X \times Y \rightarrow \mathbb{R}$ is a continuous function that is concave-convex, i.e.

$f(\cdot, y) : X \rightarrow \mathbb{R}$ is [concave](#) for fixed y , and

$f(x, \cdot) : Y \rightarrow \mathbb{R}$ is [convex](#) for fixed x .

Then we have that

$$\max_{x \in X} \min_{y \in Y} f(x, y) = \min_{y \in Y} \max_{x \in X} f(x, y).$$

Sion's minimax theorem

From Wikipedia, the free encyclopedia

In [mathematics](#), and in particular [game theory](#), **Sion's minimax theorem** is a generalization of [John von Neumann's minimax theorem](#), named after [Maurice Sion](#).

It states:

Let X be a [compact convex](#) subset of a [linear topological space](#) and Y a convex subset of a linear topological space. If f is a real-valued [function](#) on $X \times Y$ with

$f(x, \cdot)$ [upper semicontinuous](#) and [quasi-concave](#) on Y , $\forall x \in X$, and

$f(\cdot, y)$ lower semicontinuous and quasi-convex on X , $\forall y \in Y$

then,

$$\min_{x \in X} \sup_{y \in Y} f(x, y) = \sup_{y \in Y} \min_{x \in X} f(x, y).$$

Notation


- **Last lectures** (binary) classifier outputs 0/1. In this case the Bayes classifier has form $\mathbb{I}\{\mathbb{E}[Y|X] > 1/2\}$.
- **This lecture** (binary) classifier outputs -1/1. In this case the Bayes classifier has form $\mathbb{I}\{\mathbb{E}[Y|X] > 0\} = \text{sign}(\mathbb{E}[Y|X])$.

Usually classifiers have form $h(x) = \text{sign}(H(x))$. Examples include classification based on logistic regression, k -nearest-neighbors, boosting.


Decision stump: $h(X_i) = 2\mathbf{1}(e_j^T X_i \geq c) - 1$, for some j, c

Decision list: a sequence of if/else decision stumps

Decision tree: a tree of if/else decision stumps



(a) Decision stump performs well.



(b) Decision stump fails. However, decision lists does well

Algorithm 1 Decision list example

```

if  $e_1^\top x_i > 3.5$  then
    Predict +1
else if  $e_2^\top x_i > 3.5$  then
    Predict +1
else
    Predict -1
end if
  
```




Figure 5.2: Decision list perform poorly. However, decision tree performs well

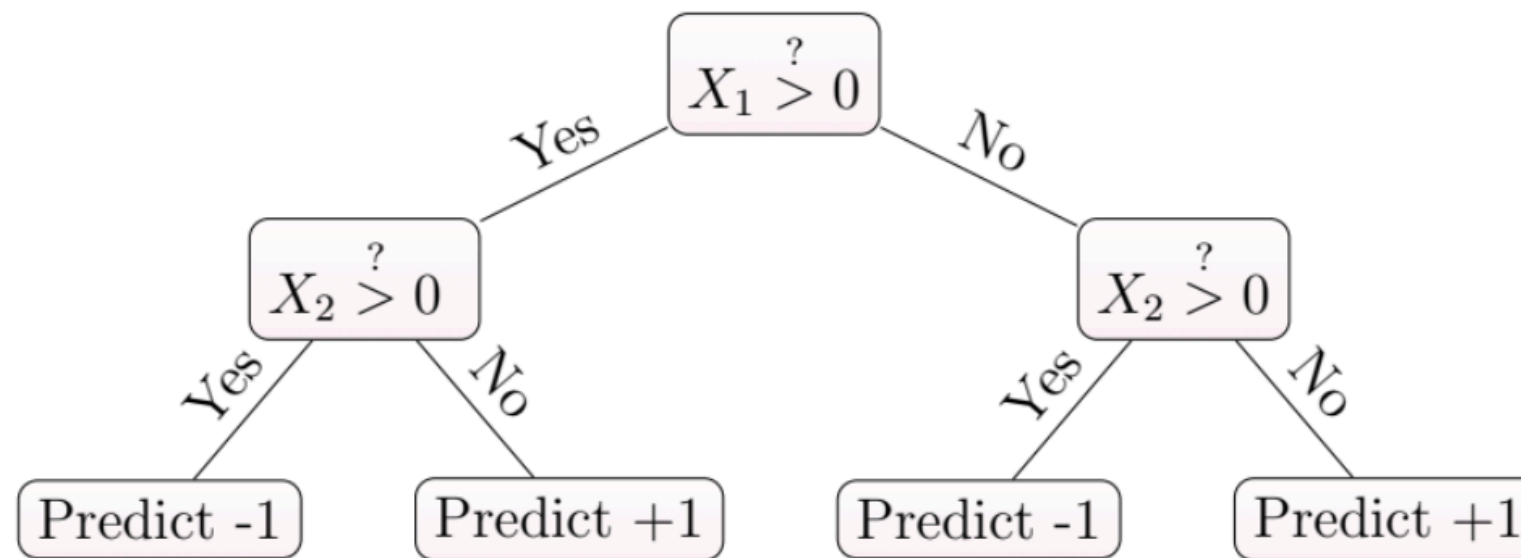


Figure 5.3: Example of a decision tree

Edge or “margin”, weighted edge

“Edge” of a classifier Edge of classifier $h \in \mathcal{H}$ is defined as $\frac{1}{n} \sum_{i=1}^n Y_i h(X_i)$. Edge provides a way to describe how much better than chance the classifier is:

- Assume that there is a perfect classifier $h^* : h^*(X_i) = Y_i, \forall i \in \{1, \dots, n\}$. Then its edge is simply equal to 1.
- Consider, on the contrary, a random guess classifier. It is trivial to show that with high probability its edge concentrates around 0.

Error of a classifier in this case can be viewed as: $\frac{1}{2} \cdot (1 - \text{edge})$. For further analysis we define a “weighted” edge as $\sum_{i=1}^n w_i Y_i h(X_i)$ where weights satisfy:

$$\sum_{i=1}^n w_i = 1, \quad w_i > 0, \quad \forall i \in \{1, \dots, n\}$$

In the previous definition each data point is equally weighted with weight $1/n$.

Weak learning hypothesis

Weak learning hypothesis: $\exists \gamma > 0$, such that for any set of weights w , there is a classifier $h \in \mathcal{H}$ whose weighted edge is at least γ .

$$\text{Let } M(r, c) = h_r(X_c)Y_c \quad \forall w \in \Delta_C \exists r \in [R] \quad e_r^T M w \geq \gamma$$

Strong learning: \exists a classifier in $\text{span}(\mathcal{H})$ with zero training error

$$\exists p \in \Delta_R \quad \forall q \in \Delta_C \quad p^T M q \geq \gamma$$

$$\exists p \in \Delta_R \quad \forall c \in [C] \quad p^T M e_c \geq \gamma$$

every element of $p^T M$ is positive

$f(X_c) := \text{sign}(p^T M e_c)$ has zero training error, for some $p \in \Delta_R$

The breakthrough Weak learning implies strong learning!

But how do we find this mixture $p \in \Delta_R$ of classifiers?

Outline

1. Weak learning implies strong learning ($1/2$ class)

2. *Adaboost* ($1/2$ class)

Algorithm 1 AdaBoost algorithm

for $m = 1, \dots, M$ **do**

(1) Compute weighted error:

$$\varepsilon(h) = \sum_{i=1}^n w_i \mathbb{I}\{Y_i \neq h(X_i)\}$$

Find a classifier h_m :

$$h_m = \arg \min_{h \in \mathcal{H}} \varepsilon(h)$$

or pick any h with nontrivial edge

(2) Compute:

$$\alpha_m = \frac{1}{2} \log \left(\frac{1 - \varepsilon_m}{\varepsilon_m} \right)$$

$$\epsilon_m = \epsilon(h_m)$$

(3) Update weights as:

$$w_i \leftarrow \frac{w_i e^{-\alpha_m Y_i h_m(X_i)}}{Z_m}$$

where Z is a normalization

end for

Output the classifier:

$$\begin{aligned} Z_m &= \sum_{i=1}^n w_m(i) \exp(-\alpha_m y_i h(x_i)) \\ &= \sum_{i: y_i h_m(x)=1} w_m(i) \exp(-\alpha_m) + \sum_{i: y_i h_m(x)=-1} w_m(i) \exp(\alpha_m) \\ &= (1 - \varepsilon_m) \exp(-\alpha_m) + \varepsilon_m \exp(\alpha_m) \\ &= 2\sqrt{\varepsilon_m(1 - \varepsilon_m)} \end{aligned}$$




Figure 5.4: This figure is taken from [2]: each learning curve shows the training and test error curves (lower and upper curves, respectively) of the combined classifier as a function of the number of classifiers combined.

Theorem 2 (Convergence of empirical risk of Adaboost) *The empirical risk of the output of Adaboost algorithm 1 $\hat{R}(f)$ satisfies:*

$$\begin{aligned} \hat{R}(f) &\leq \exp\left(-2 \sum_{m=1}^M \left(\frac{1}{2} - \varepsilon_m\right)^2\right) \\ &\leq \exp(-2M\gamma^2) \text{ if weak learning hypothesis is true} \end{aligned} \tag{6.3}$$

If $M > \frac{\log n}{2\gamma^2}$, then $\hat{R}(f) < 1/n$, and hence $\hat{R}(f) = 0$

Theorem 2 (Convergence of empirical risk of Adaboost) *The empirical risk of the output of Adaboost algorithm 1 $\hat{R}(f)$ satisfies:*

$$\hat{R}(f) \leq \exp(-2 \sum_{m=1}^M (\frac{1}{2} - \varepsilon_m)^2) \quad (6.3)$$

$$\begin{aligned} \hat{R}(f) &= \frac{1}{n} \sum_{i=1}^n \mathbb{I}(y_i \neq f(x_i)) \\ &\leq \frac{1}{n} \sum_{i=1}^n \exp(-y_i f(x_i)) \\ &= \frac{1}{n} \sum_{i=1}^n [n \prod_{m=1}^M Z_m] w_{M+1}(i) \\ &= \prod_{m=1}^M Z_m \end{aligned}$$

Surrogate loss

$$w_{M+1}(i) = \frac{w_M(i) e^{-\alpha_M Y_i h_M(X_i)}}{Z_M} = \frac{e^{-Y_i \sum_m \alpha_m h_m(X_i)}}{n \prod_{m=1}^M Z_m}$$

α_m minimizes weighted loss

$$\begin{aligned} Z_m &= \sum_{i=1}^n w_m(i) \exp(-\alpha_m y_i h(x_i)) \\ &= \sum_{i: y_i h_m(x)=1} w_m(i) \exp(-\alpha_m) + \sum_{i: y_i h_m(x)=-1} w_m(i) \exp(\alpha_m) \\ &= (1 - \varepsilon_m) \exp(-\alpha_m) + \varepsilon_m \exp(\alpha_m) \\ &= 2 \sqrt{\varepsilon_m (1 - \varepsilon_m)} \leq \exp(-2(\frac{1}{2} - \varepsilon_m)^2) \end{aligned}$$

Convex surrogate loss minimization by coordinate descent

$$\hat{R}(\beta) = \frac{1}{n} \sum_{i=1}^n \exp \left(-y_i \sum_{j=1}^J \beta_j h_j(x_i) \right)$$

Adaboost solves $\min_{\beta \in \mathbb{R}_+^J} \hat{R}(\beta) = \min_{f \in \text{span}(\mathcal{H})} \hat{R}(f)$ by “coordinate descent”.

1. Begin at $\beta^{(0)} = [0, 0, \dots, 0]$
2. At step t , pick direction $e_t \in \{e_j\}_{j \in J}$ and stepsize $\alpha_t \geq 0$ to minimize $\hat{R}(\beta^{(t-1)} + \alpha_t e_t)$
3. Gradient with respect to coordinate j is $\hat{R}'(\beta^{t-1})_j \propto (2\epsilon_{t,j} - 1) \prod_{s=1}^{t-1} Z_s$, where $\epsilon_{t,j}$ is weighted error of h_j
4. h_t is chosen to minimize the weighted error, optimal stepsize happens to equal $\log\left(\frac{1 - \epsilon_t}{\epsilon_t}\right)$

