

# Results of the survey

## +Von Neumann's minimax theorem

Aaditya Ramdas

Dept. of Statistics and Data Science  
Machine Learning Dept.  
Carnegie Mellon University

# Outline

1. Weak learning implies strong learning ( $1/2$  class)

**2. *Adaboost* ( $1/2$  class)**

# Recap

A “mixed” strategy is a distribution over actions.

Expected payoff is  $\mathbb{E}_{r \sim p, c \sim q}[M(r, c)] = p^T M q = \sum_{r \in [R]} \sum_{c \in C} M(r, c) p_r q_c$

**Theorem:**  $\min_{q \in \Delta_C} \max_{p \in \Delta_R} p^T M q = \max_{p \in \Delta_R} \min_{q \in \Delta_C} p^T M q = v^*$

Value of the game

$e_j$  is the canonical basis vector  $[0, \dots, 1, \dots, 0]$

**Implications:**

$$\begin{aligned} \exists p \in \Delta_R \forall q \in \Delta_C \quad p^T M q &\geq v^* \\ \forall q \in \Delta_C \exists i \in [R] \quad e_i^T M q &\geq v^* \\ \exists q \in \Delta_C \forall p \in \Delta_R \quad p^T M q &\leq v^* \\ \forall p \in \Delta_R \exists j \in [C] \quad p^T M e_j &\leq v^* \end{aligned}$$

**For “mixed strategies”, order does not matter!**

## Zero-sum games [[edit](#)]

---

The minimax theorem was first proven and published in 1928 by [John von Neumann](#),<sup>[3]</sup> who is quoted as saying "*As far as I can see, there could be no theory of games ... without that theorem ... I thought there was nothing worth publishing until the Minimax Theorem was proved*".<sup>[4]</sup>

Formally, von Neumann's minimax theorem states:

Let  $X \subset \mathbb{R}^n$  and  $Y \subset \mathbb{R}^m$  be [compact convex](#) sets. If  $f : X \times Y \rightarrow \mathbb{R}$  is a continuous function that is concave-convex, i.e.

$f(\cdot, y) : X \rightarrow \mathbb{R}$  is [concave](#) for fixed  $y$ , and

$f(x, \cdot) : Y \rightarrow \mathbb{R}$  is [convex](#) for fixed  $x$ .

Then we have that

$$\max_{x \in X} \min_{y \in Y} f(x, y) = \min_{y \in Y} \max_{x \in X} f(x, y).$$

## Sion's minimax theorem

---

From Wikipedia, the free encyclopedia

In [mathematics](#), and in particular [game theory](#), **Sion's minimax theorem** is a generalization of [John von Neumann's minimax theorem](#), named after [Maurice Sion](#).

It states:

Let  $X$  be a [compact convex](#) subset of a [linear topological space](#) and  $Y$  a convex subset of a linear topological space. If  $f$  is a real-valued [function](#) on  $X \times Y$  with

$f(x, \cdot)$  [upper semicontinuous](#) and [quasi-concave](#) on  $Y$ ,  $\forall x \in X$ , and

$f(\cdot, y)$  lower semicontinuous and quasi-convex on  $X$ ,  $\forall y \in Y$

then,

$$\min_{x \in X} \sup_{y \in Y} f(x, y) = \sup_{y \in Y} \min_{x \in X} f(x, y).$$

# Notation


- **Last lectures** (binary) classifier outputs 0/1. In this case the Bayes classifier has form  $\mathbb{I}\{\mathbb{E}[Y|X] > 1/2\}$ .
- **This lecture** (binary) classifier outputs -1/1. In this case the Bayes classifier has form  $\mathbb{I}\{\mathbb{E}[Y|X] > 0\} = \text{sign}(\mathbb{E}[Y|X])$ .

Usually classifiers have form  $h(x) = \text{sign}(H(x))$ . Examples include classification based on logistic regression,  $k$ -nearest-neighbors, boosting.


Decision stump:  $h(X_i) = 2\mathbf{1}(e_j^T X_i \geq c) - 1$ , for some  $j, c$

Decision list: a sequence of if/else decision stumps

Decision tree: a tree of if/else decision stumps



(a) Decision stump performs well.



(b) Decision stump fails. However, decision lists does well

---

**Algorithm 1** Decision list example

---

```

if  $e_1^\top x_i > 3.5$  then
    Predict +1
else if  $e_2^\top x_i > 3.5$  then
    Predict +1
else
    Predict -1
end if
  
```

---




Figure 5.2: Decision list perform poorly. However, decision tree performs well

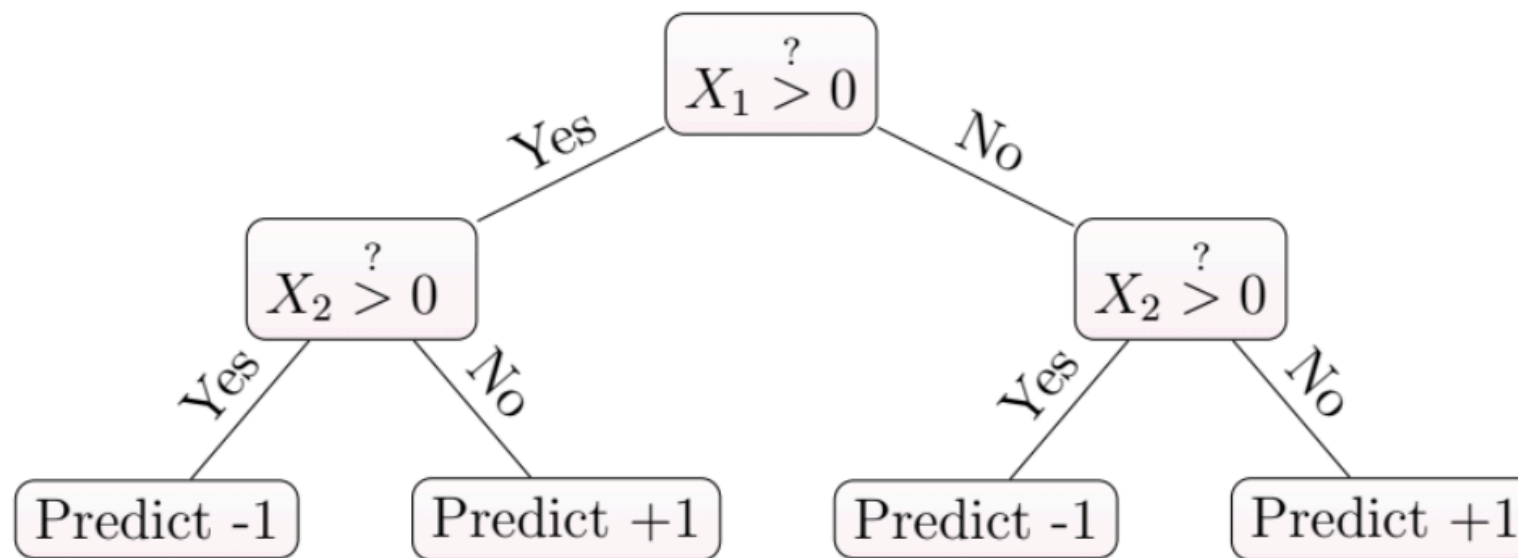


Figure 5.3: Example of a decision tree

# Edge or “margin”, weighted edge

**“Edge” of a classifier** Edge of classifier  $h \in \mathcal{H}$  is defined as  $\frac{1}{n} \sum_{i=1}^n Y_i h(X_i)$ . Edge provides a way to describe how much better than chance the classifier is:

- Assume that there is a perfect classifier  $h^* : h^*(X_i) = Y_i, \forall i \in \{1, \dots, n\}$ . Then its edge is simply equal to 1.
- Consider, on the contrary, a random guess classifier. It is trivial to show that with high probability its edge concentrates around 0.

Error of a classifier in this case can be viewed as:  $\frac{1}{2} \cdot (1 - \text{edge})$ . For further analysis we define a “weighted” edge as  $\sum_{i=1}^n w_i Y_i h(X_i)$  where weights satisfy:

$$\sum_{i=1}^n w_i = 1, \quad w_i > 0, \quad \forall i \in \{1, \dots, n\}$$

In the previous definition each data point is equally weighted with weight  $1/n$ .



# Weak learning hypothesis

Weak learning hypothesis:  $\exists \gamma > 0$ , such that for any set of weights  $w$ , there is a classifier  $h \in \mathcal{H}$  whose weighted edge is at least  $\gamma$ .

Let  $M(r, c) = h_r(X_c)Y_c$        $\forall w \in \Delta_C \exists r \in [R] \quad e_r^T M w \geq \gamma$

Strong learning:  $\exists$  a classifier in  $\text{span}(\mathcal{H})$  with zero training error

$$\exists p \in \Delta_R \quad \forall q \in \Delta_C \quad p^T M q \geq \gamma$$

$$\exists p \in \Delta_R \quad \forall c \in [C] \quad p^T M e_c \geq \gamma$$

every element of  $p^T M$  is positive

$h(X_c) := \text{sign}(p^T M e_c)$  has zero training error, for some  $p \in \Delta_R$

**The breakthrough**      Weak learning implies strong learning!

But how do we find this mixture  $p \in \Delta_R$  of classifiers?

# Outline

1. Weak learning implies strong learning ( $1/2$  class)

**2. *Adaboost* ( $1/2$  class)**

---

**Algorithm 1** AdaBoost algorithm

---

**for**  $m = 1, \dots, M$  **do**

(1) Compute weighted error:

$$\varepsilon(h) = \sum_{i=1}^n w_i \mathbb{I}\{Y_i \neq h(X_i)\}$$

Find a classifier  $h_m$ :

$$h_m = \arg \min_{h \in \mathcal{H}} \varepsilon(h)$$

or pick any  $h$  with nontrivial edge

(2) Compute:

$$\alpha_m = \frac{1}{2} \log \left( \frac{1 - \varepsilon_m}{\varepsilon_m} \right)$$

$$\epsilon_m = \epsilon(h_m)$$

(3) Update weights as:

$$w_i \leftarrow \frac{w_i e^{-\alpha_m Y_i h_m(X_i)}}{Z_m}$$

where  $Z$  is a normalization

**end for**

Output the classifier:

$$\begin{aligned} Z_m &= \sum_{i=1}^n w_m(i) \exp(-\alpha_m y_i h(x_i)) \\ &= \sum_{i: y_i h_m(x)=1} w_m(i) \exp(-\alpha_m) + \sum_{i: y_i h_m(x)=-1} w_m(i) \exp(\alpha_m) \\ &= (1 - \varepsilon_m) \exp(-\alpha_m) + \varepsilon_m \exp(\alpha_m) \\ &= 2\sqrt{\varepsilon_m(1 - \varepsilon_m)} \end{aligned}$$

---




Figure 5.4: This figure is taken from [2]: each learning curve shows the training and test error curves (lower and upper curves, respectively) of the combined classifier as a function of the number of classifiers combined.

**Theorem 2 (Convergence of empirical risk of Adaboost)** *The empirical risk of the output of Adaboost algorithm 1  $\hat{R}(f)$  satisfies:*

$$\begin{aligned} \hat{R}(f) &\leq \exp\left(-2 \sum_{m=1}^M \left(\frac{1}{2} - \varepsilon_m\right)^2\right) \\ &\leq \exp(-2M\gamma^2) \text{ if weak learning hypothesis is true} \end{aligned} \tag{6.3}$$

If  $M > \frac{\log n}{2\gamma^2}$ , then  $\hat{R}(f) < 1/n$ , and hence  $\hat{R}(f) = 0$

**Theorem 2 (Convergence of empirical risk of Adaboost)** *The empirical risk of the output of Adaboost algorithm 1  $\hat{R}(f)$  satisfies:*

$$\hat{R}(f) \leq \exp(-2 \sum_{m=1}^M (\frac{1}{2} - \varepsilon_m)^2) \quad (6.3)$$

$$\begin{aligned} \hat{R}(f) &= \frac{1}{n} \sum_{i=1}^n \mathbb{I}(y_i \neq f(x_i)) \\ &\leq \frac{1}{n} \sum_{i=1}^n \exp(-y_i f(x_i)) \\ &= \frac{1}{n} \sum_{i=1}^n \left[ n \prod_{m=1}^M Z_m \right] w_{M+1}(i) \\ &= \prod_{m=1}^M Z_m \end{aligned}$$

Surrogate loss

$$w_{M+1}(i) = \frac{w_M(i) e^{-\alpha_M Y_i h_M(X_i)}}{Z_M} = \frac{e^{-Y_i \sum_m \alpha_m h_m(X_i)}}{n \prod_{m=1}^M Z_m}$$

$\alpha_m$  minimizes weighted loss

$$\begin{aligned} Z_m &= \sum_{i=1}^n w_m(i) \exp(-\alpha_m y_i h(x_i)) \\ &= \sum_{i: y_i h_m(x)=1} w_m(i) \exp(-\alpha_m) + \sum_{i: y_i h_m(x)=-1} w_m(i) \exp(\alpha_m) \\ &= (1 - \varepsilon_m) \exp(-\alpha_m) + \varepsilon_m \exp(\alpha_m) \\ &= 2 \sqrt{\varepsilon_m (1 - \varepsilon_m)} \leq \exp(-2(\frac{1}{2} - \varepsilon_m)^2) \end{aligned}$$

# Convex surrogate loss minimization by coordinate descent

$$\hat{R}(\beta) = \frac{1}{n} \sum_{i=1}^n \exp \left( -y_i \sum_{j=1}^J \beta_j h_j(x_i) \right)$$

Adaboost solves  $\min_{\beta \in \mathbb{R}_+^J} \hat{R}(\beta) = \min_{f \in \text{span}(\mathcal{H})} \hat{R}(f)$  by “coordinate descent”.

1. Begin at  $\beta^{(0)} = [0, 0, \dots, 0]$
2. At step  $t$ , pick direction  $e_t \in \{e_j\}_{j \in J}$  and stepsize  $\alpha_t \geq 0$  to minimize  $\hat{R}(\beta^{(t-1)} + \alpha_t e_t)$
3. Gradient with respect to coordinate  $j$  is  $\hat{R}'(\beta^{t-1})_j \propto (2\epsilon_{t,j} - 1) \prod_{s=1}^{t-1} Z_s$ , where  $\epsilon_{t,j}$  is weighted error of  $h_j$
4.  $h_t$  is chosen to minimize the weighted error, optimal stepsize happens to equal  $\log\left(\frac{1 - \epsilon_t}{\epsilon_t}\right)$

