

t-Distributed Stochastic Neighbor Embedding (t-SNE)

Slides: Selina Carter


April 19, 2022

Outline

1. Motivation
2. What is “manifold learning”
3. SNE algorithm
+ Examples
4. t-SNE algorithm
+ Examples

Motivation




- PCA maps data in high dimensions to a plane



Graphic: [Matthias Sholz](#)

Motivation


- But what if your data isn't nicely represented by a plane?




Graphics: [Wolfram](#), [Gashler et al 2011](#)

This is called “manifold learning”


- **Manifold:** basically, a type of curve or surface (locally resembles a Euclidean space)
 - In our case, we’re talking about k-manifolds that are subspaces of a Euclidean space
- Simple example: 2-manifold (a surface)




- Not manifolds:



Three or more faces share an edge.



Two or more faces share a single vertex but no edge.



Adjacent faces have opposite normals.

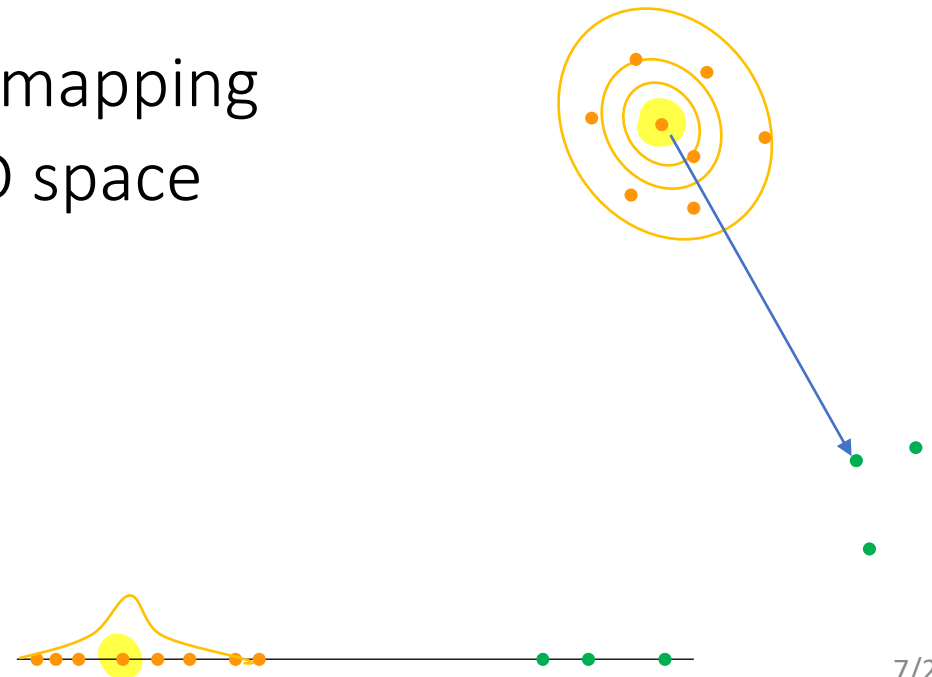
Graphics: [Wikipedia](#), [Autodesk](#)

Manifold learning is not new

- Linear manifolds (subspaces)
 - Principal Component Analysis (PCA) (Pearson, 1901; Hotelling, 1933)
 - (Classical) Multidimensional Scaling (MDS) (Torgerson, 1952)
- Non-linear manifolds
 - Sammon mapping (Sammon, 1969)
 - Curvilinear components analysis (CCA) (Demartines and Herault, 1997)
 - Isomap (Tenenbaum et al., 2000)
 - Locally Linear Embedding (LLE) (Roweis and Saul, 2000)
 - Stochastic Neighbor Embedding (SNE) (Hinton and Roweis, 2002)
 - Laplacian Eigenmaps (Belkin and Niyogi, 2002)
 - Maximum Variance Unfolding (MVU) (Weinberger et al., 2004)
 - t-Distributed Stochastic Neighbor Embedding (t-SNE) (Van der Maaten and Hinton, 2008)

Stochastic neighbor embedding (SNE)

- Big picture:
 - **Dimensionality reduction**: describe distances between objects in a high-D space by placing the objects in a low-D space
 - “**embed**” objects originally in high-D space into a lower-D space
- Toolkit:
 - Probabilistic approach: there is no “one” mapping
 - Gaussian distribution in high-D and low-D space
 - Gradient descent



SNE algorithm (part 1)

1. Define a “distance metric” between points i and j in original (high-dimensional) space ($\mathbf{x}_i, \mathbf{x}_j \in \mathbb{R}^N$):

$$d_{ij}^2 := \frac{\|\mathbf{x}_i - \mathbf{x}_j\|^2}{2\sigma_i^2}, \quad \mathbf{x}_i, \mathbf{x}_j \in \mathbb{R}^N, \sigma_i \in \mathbb{R}$$

Note:


$$d_{ij} \neq d_{ji}$$

2. For each object i and each potential neighbor j , compute the “similarity of \mathbf{x}_j to \mathbf{x}_i ”:
i.e., conditional probability $p_{j|i}$ that i picks j as its neighbor:

$$p_{j|i} := \frac{\exp(-d_{ij}^2)}{\sum_{k \neq i} \exp(-d_{ik}^2)}, \quad p_{i|i} = 0$$

Note:

$$p_{j|i} \neq p_{i|j}$$



SNE algorithm (part 2)

3. Now in low-dimensions ($\mathbb{R}^m, m \ll N$): imagine low-dimensional images $\mathbf{y}_i \in \mathbb{R}^m$ of each point i .

The conditional probability $q_{j|i}$ that i picks j as its neighbor is:

$$q_{j|i} := \frac{\exp\left(-\|\mathbf{y}_i - \mathbf{y}_j\|^2\right)}{\sum_{k \neq i} \exp\left(-\|\mathbf{y}_i - \mathbf{y}_k\|^2\right)}, \quad q_{i|i} = 0$$

Note:

$q_{j|i} \neq q_{i|j}$

4. **Goal:** choose mappings $\{\mathbf{y}_1, \dots, \mathbf{y}_n\}$ so that $p_{j|i} = q_{j|i}$

➡ For each object i , minimize cost function: sum of Kullback-Leibler divergences between $p_{j|i}$ and $q_{j|i}$:

$$C := \sum_{i=1}^N \sum_{j=1}^N p_{j|i} \log\left(\frac{p_{j|i}}{q_{j|i}}\right) = \sum_{i=1}^N KL(P_i || Q_i)$$

Look closer at the cost function

$$C := \sum_{i=1}^N \sum_{j=1}^N p_{j|i} \log \left(\frac{p_{j|i}}{q_{j|i}} \right) = \sum_{i=1}^N KL(P_i || Q_i)$$

- With respect to *what* are we minimizing?

$$\min_{\mathbf{y}_1, \dots, \mathbf{y}_{n \in \mathbb{R}^m}} C = \min_{\mathbf{y}_1, \dots, \mathbf{y}_{n \in \mathbb{R}^m}} \sum_{i=1}^N KL(P_i || Q_i)$$

- How do you minimize it?
 - Naive: minimize ~~$\frac{p_{j|i}}{q_{j|i}}$~~
 - You have a “budget” for $q_{j|i}$: $\sum_{j=1}^N q_{j|i} = 1$
 - Result: choose $\mathbf{y}_1, \dots, \mathbf{y}_{n \in \mathbb{R}^m}$ such that $q_{j|i} \approx p_{j|i}$
- Note: KL divergence is not symmetric in p and q
 - Cost(small $q_{j|i}$, large $p_{j|i}$) $>$ Cost(large $q_{j|i}$, small $p_{j|i}$)

SNE algorithm (review)

- Fixed things:

1. Data: $\mathbf{x}_1, \dots, \mathbf{x}_n \in \mathbb{R}^N$
2. Choose size of lower dimension ($\mathbb{R}^m, m \ll N$)
3. Distance metric d_{ij}^2 in \mathbb{R}^N

$$d_{ij}^2 := \frac{\|\mathbf{x}_i - \mathbf{x}_j\|^2}{2\sigma_i^2}, \quad \sigma_i \in \mathbb{R}$$

4. In \mathbb{R}^N , compute probability i picks j as its neighbor:

$$p_{j|i} := \frac{\exp(-d_{ij}^2)}{\sum_{k \neq i} \exp(-d_{ik}^2)}$$

5. Given the images $\mathbf{y}_1, \dots, \mathbf{y}_n \in \mathbb{R}^m$,

$$q_{j|i} := \frac{\exp(-\|\mathbf{y}_i - \mathbf{y}_j\|^2)}{\sum_{k \neq i} \exp(-\|\mathbf{y}_i - \mathbf{y}_k\|^2)}$$

- Task: choose $\mathbf{y}_1, \dots, \mathbf{y}_n \in \mathbb{R}^m$ such that $q_{ij} \approx p_{ij}$


➡ Use gradient descent: $\mathbf{y}^{(t)} = \mathbf{y}^{(t-1)} + \eta \frac{\delta \mathcal{C}}{\delta \mathbf{y}} + \alpha(t)(\mathbf{y}^{(t-1)} - \mathbf{y}^{(t-2)})$

Example 1: SNE


- Horseshoe shape:

1. Data: $\mathbf{x}_1, \dots, \mathbf{x}_{500} \in \mathbb{R}^{10}$
2. Lower dimension: \mathbb{R}^1 and \mathbb{R}^2


Original data
(horseshoe jittered in \mathbb{R}^{10})



SNE to \mathbb{R}^1




SNE to \mathbb{R}^2




Example 2: SNE

- Hand-written digits:
 1. Data: $\mathbf{x}_1, \dots, \mathbf{x}_{3000} \in \mathbb{R}^{256}$
(grayscale images)
 2. Lower dimension: \mathbb{R}^2
- Problems:
 - Difficult to optimize the cost function
 - “Crowding” of points



Graphic: [Hinton and Roweis 2002](#)

Stochastic neighbor embedding (t-SNE)

- Big picture:
 - Same idea as SNE, but easier to optimize – and “works better.”
- Key differences from SNE:
 - New cost function
 - Symmetric probabilities
 - Simpler gradient

➡ Faster to optimize
 - Instead of Gaussian, uses Student-t distribution ($df = 1$; same as Cauchy distribution) to compute similarity of points in low-D space
 - ➡ Fixes “crowding” by separating clusters

t-SNE algorithm (part 1) (changes to SNE are in yellow)

1. Define a distance metric between points i and j in original (high-dimensional) space ($\mathbf{x}_i, \mathbf{x}_j \in \mathbb{R}^N$):

$$d_{ij}^2 := \frac{\|\mathbf{x}_i - \mathbf{x}_j\|^2}{2\sigma_i^2}, \quad \mathbf{x}_i, \mathbf{x}_j \in \mathbb{R}^N, \sigma_i \in \mathbb{R}$$

2. For each object i and each potential neighbor j , compute probability p_{ij} that i picks j as its neighbor:

$$p_{ij} := \frac{p_{j|i} + p_{i|j}}{2n}, \quad p_{ii} = 0$$

Note:
 $p_{ij} = p_{ji}$

t-SNE algorithm (part 2) (changes to SNE are in yellow)

3. Now in low-dimensions ($\mathbb{R}^m, m \ll N$): imagine low-dimensional images $\mathbf{y}_i \in \mathbb{R}^m$ of each point i .

Then the induced probability q_{ij} that i picks j as its neighbor is:

$$q_{ij} = \frac{\left(1 + \|\mathbf{y}_i - \mathbf{y}_j\|^2\right)^{-1}}{\sum_{k \neq l} \exp\left(1 + \|\mathbf{y}_k - \mathbf{y}_l\|^2\right)^{-1}}$$

Note:

$$q_{ij} = q_{ji}$$

← Fixed normalizing constant across all points

4. **Goal:** choose mappings $\{\mathbf{y}_1, \dots, \mathbf{y}_n\}$ so that $p_{j|i} = q_{j|i}$

For each object i , minimize cost function: sum of Kullback-Leibler divergences between p_{ij} and q_{ij} :

$$C = \sum_{i=1}^N \sum_{j=1}^N p_{ij} \log \left(\frac{p_{ij}}{q_{ij}} \right) = KL(P||Q)$$

← No longer a sum

t-SNE algorithm (condensed)

Algorithm 1: Simple version of t-Distributed Stochastic Neighbor Embedding.

Data: data set $\mathcal{X} = \{x_1, x_2, \dots, x_n\}$,

cost function parameters: perplexity $Perp$,

optimization parameters: number of iterations T , learning rate η , momentum $\alpha(t)$.

Result: low-dimensional data representation $\mathcal{Y}^{(T)} = \{y_1, y_2, \dots, y_n\}$.

begin

 compute pairwise affinities $p_{j|i}$ with perplexity $Perp$ (using Equation 1)

 set $p_{ij} = \frac{p_{j|i} + p_{i|j}}{2n}$

 sample initial solution $\mathcal{Y}^{(0)} = \{y_1, y_2, \dots, y_n\}$ from $\mathcal{N}(0, 10^{-4}I)$

for $t=1$ **to** T **do**

 compute low-dimensional affinities q_{ij} (using Equation 4)

 compute gradient $\frac{\delta C}{\delta \mathcal{Y}}$ (using Equation 5)


 set $\mathcal{Y}^{(t)} = \mathcal{Y}^{(t-1)} + \eta \frac{\delta C}{\delta \mathcal{Y}} + \alpha(t) (\mathcal{Y}^{(t-1)} - \mathcal{Y}^{(t-2)})$

end

end

Example 2: t-SNE



- Hand-written digits (MNIST):
 1. Data: $\mathbf{x}_1, \dots, \mathbf{x}_{6000} \in \mathbb{R}^{784}$
(grayscale images)
 2. Lower dimension: \mathbb{R}^2
 3. Then used k=20 nearest neighbors (color version)



Graphic: [Hinton and Roweis 2002](#)

Example 3: PCA vs. t-SNE

- Hand-written digits (MNIST):
 1. Data: $\mathbf{x}_1, \dots, \mathbf{x}_{6000} \in \mathbb{R}^{784}$
(grayscale images)
 2. Lower dimension: \mathbb{R}^2



Example 4: t-SNE

- Facial Expression Recognition (FER)
Japanese Female Facial Expression (JAFFE) by Jizheng Yi et al, 2013




TABLE I. THE FER RATES BETWEEN DIFFERENT ALGORITHMS

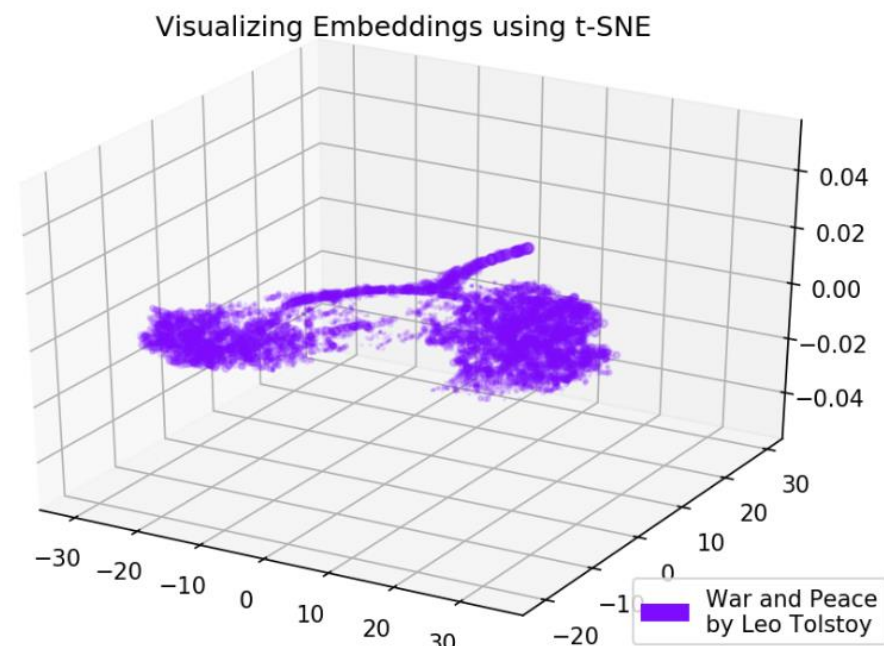
	PCA	LDA	LLE	SNE	t-SNE
SVM	73.5%	74.3%	84.7%	89.6%	90.3%
AdaboostM2	75.4%	75.9%	87.7%	90.6%	94.5%

Example 5: t-SNE

- Word embeddings
 - Words in Google News (100 billion words, \mathbb{R}^{200})
 - *War and Peace* by Tolstoy



Clusters of similar words from Google News (preplexity=15)



Visualization of the Word2Vec model trained on War and Peace

Example 6: t-SNE

- Interactive viz & tutorial: <https://distill.pub/2016/misread-tsne/>

Questions?